Statistical Learning – Data science - 2020/21 – Exercise 1 - 11/03/2021

Exercise 1: Telco Customer Churn first data analysis using Python
Please, execute the following tasks and provide answers to the proposed questions.

1. Open the Telco Customer Churn dataset page in Kaggle.
- Hint: https://www.kaggle.com/blastchar/telco-customer-churn
- Have a look to the "Overview" tab to understand something more about the dataset

2. Check the main properties of this dataset in the "Data" tab.
- How many samples (rows) does it have?
- How many variables (columns)?
- What does each row/column represent?
- Which is the "target" column? What does it represent?

3. Download the dataset into your computer.
- Which is the extension of the downloaded file?

4. Uncompress the file
- Which is the extension of the uncompressed file?

5. Open the uncompressed file by both a text editor and a spreadsheet software
- Which symbol is used to separate columns?
- Which symbol is used to separate rows?
- Which values can you find for variable SeniorCitizen? And for variable Partner?

6. Generate a new notebook for analyzing this dataset
- Hint: click on "New Kernel", then choose the Notebook kernel type, on the right
- Assign the following title to the notebook: SL_L1_TelcoCustomerChurn_YourSurname
- Then click on the "Commit" button on top-right to make the notebook ready to be started

7. Open the notebook documentation page to get help if needed
- Hint: click the "Docs" link on the right-bottom of your notebook page

8. Select the first cell (we will call it "Library import cell" in the following), run it
- What is the output of this action?
- What does the code "*import numpy as np*" do? Can you provide a reference website for this library?
- What does the code "*import pandas as pd*" do? Can you provide a reference website for this library?
- What does the code "*import os*" do? Can you provide a reference website for this library?
- How many data files are available? Please provide their names.

9. Add to the first cell new lines to load the following libraries: seaborn, matplotlib.pyplot, sklearn.linear_model (only LogisticRegression)
- Hint: find similar code in the Titanic notebook if needed

10. Select the first cell and add a new cell on top of it
- Hint: use the button on top-right of the cell

11. Select the new cell and transform it in a "Markdown" cell, then copy all the text in this pdf file and paste it in the new Markdown cell

12. Please write your answers to the questions above in the new Markdown cell. From now on you can use the same cell to write your answers as well

13. Select the "Library input cell" and add a new cell below it

14. Use the new cell to load the Telco Customer Churn dataset into a Pandas DataFrame variable called *data*

- Hint: find similar code in the Titanic notebook if needed
- Remind to run the cell after writing the code-box

15. Add the following comment before data loading line: "Data acquisition"

16. Add also a Markdown cell before the data loading cell and write in bold the text "Data acquisition"

- Markdown cells should be used to give a structure to the report, hence they should be added before each new section

17. In a new cell show the number of rows, the number of columns, and the total number of cells in the dataset

- Hint: display the related *parameters* of the Pandas DataFrame
- Hint: use the *print* function to print the results
- You should print, in particular, the following strings:
  - "The number of customers is XXXX"
  - "The number of variables is YYYY"
  - "The total number of cells is ZZZZ"
- Other hints:
  - How can you select a single element from the shape tuple?
  - How can you convert a number to string?
  - How can you concatenate two strings?
  - How can you print the final string?

18. Add the following comment at the beginning of the cell: "Dataset dimension"

19. Add a new markdown cell before this cell and write in it the title "Data Analysis"

20. In a new cell show the names of the variables in the dataset

- Hint: print the *column*'s names of variable *data*

21. In a new cell show the *first* and *last* 10 rows in the dataset

- Hint: find the correct DataFrame *methods* in the Pandas' documentation

22. In a new cell show i) the type of variable *data,* ii) the number of missing values for each variable, iii) the type of each variable, iv) the total memory used to store variable data

- Hint: all this information can be provided by a single method of DataFrame
- How many missing values are there in total?
- Which variables are categorical?
- Which variables are numerical?

23. In a new cell show the following basic statistics for all numerical variables: number of non-missing values, mean, standard deviation, minimum, maximum, median, 1st and 3rd quartiles

- Hint: all this information can be provided by a single method of DataFrame

24. In a new cell show the following basic information for all categorical variables: number of non-missing values, number of unique values, most frequent value and frequency of the most frequent value.

- Hint: all this information can be provided by the DataFrame method used in question 22, using specific arguments
- Can you see any strange value in this result?

25. In a new cell show the histograms of each numeric variable (i.e., column) in the dataset

- Hint: try to find a specific method in the DataFrame API documentation

26. In a new cell show the box-plots of each numeric variable (i.e., column) in the dataset

- Hint: try to find a specific method in the DataFrame API documentation
- Does this chart provide a good visualization? Why?
- Try to generate one box-plot for each numerical variable
- Try to put all three charts in the same figure using the subplot function

27. In a new cell show the histograms of the categorical variables in the dataset

- Hint: try to use a function from the Seaborn library which counts the number of time each element appears and makes a related bar plot
- Hint: use the subplot function to put all the charts in the same figure
- Hint: resize the figure so that to avoid overlapping and enable a clear visualization of all charts

28. In a new cell generate a new DataFrame called data1 and containing only variables gender, Partner, MonthlyCharges, Churn

- Hint: you could try also other selections

29. In a new cell show the first 5 rows of the new dataset

30. Convert categorical values in data1 to numeric as follows:

- **gender: Male=0, Female=1**
- **Partner: No=0, Yes=1**
- **Churn: No=0, Yes=1**

- Hint: find similar code in the Titanic notebook if needed

31. Generate a separate Series variable called data1Churn for the dependent (churn) variable and drop it from DataFrame data1

- Hint: Series is a data structure defined in Pandas, try to find its documentation page
- Hint: each column of a DataFrame is a Series
- Hint: learn how to drop columns from a dataset in the Titanoc notebook
- What is the difference between data1[[‘Churn’]] and data1[‘Churn’]?
- When single square brackets are used with Pandas DataFrame? When double brackets are used instead?

32. Generate a linear logistic model using data1 as dependent variables and data1Churn as independent variable, then show the model “score”

- Hint: try to find a function for linear logistic model learning in the sklearn library

- Hint: find similar code in the Titanic notebook if needed

33. Show the parameters of the linear logistic model computed above. Which variable seems to be more related to customer churn?

- Hint: find similar code in the Titanic notebook if needed

34. If you want, click on the *Sharing* field on the right hand side menu and share the notebook with me (Kaggle user: albertocastellini)

- No score/evaluation will be given, don’t worry :-)